


Introdução a *Multi-threading* com C++ 11

(C++ 11 == ISO/IEC 14882:2011)

Fabio Galuppo, M.Sc.

<http://fabiogaluppo.com>

fabiogaluppo@acm.org



First year awarded:
2002

Number of MVP Awards:
11

Technical Expertise:
Visual C++

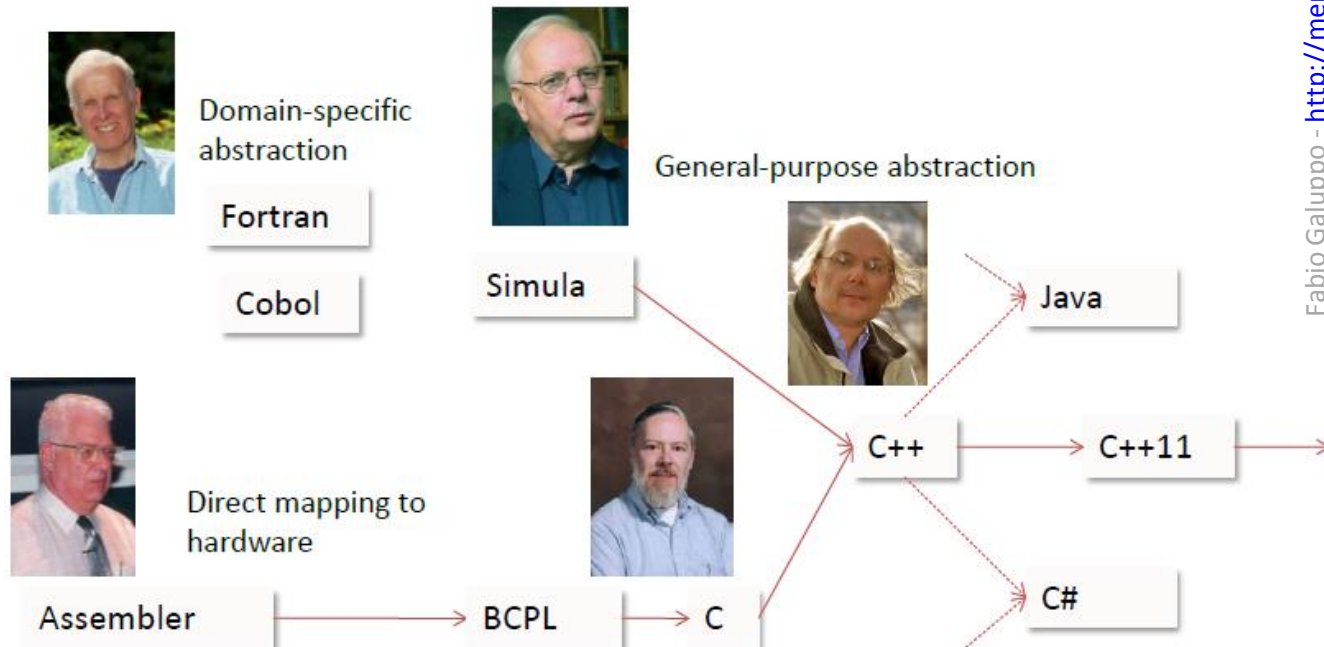

Technical Interests:
Visual C#, Visual F#

The C++ Programming Language


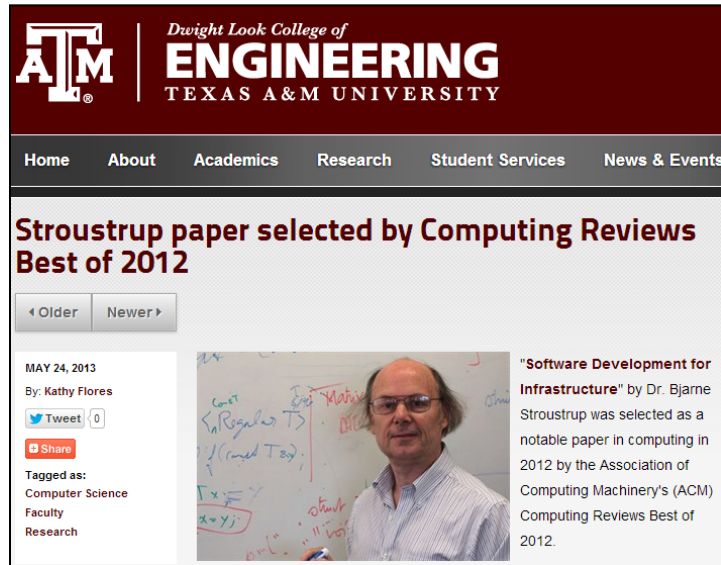
C++ is a general purpose programming language with a bias towards systems programming that

- is a better C
- supports data abstraction
- supports object-oriented programming
- supports generic programming.

C++ is a general purpose programming language designed to make programming more enjoyable for the serious programmer.



C++



IEEE Computer Magazine Volume:45, Issue: 1

<http://dx.doi.org/10.1109/MC.2011.353>

“A light-weight abstraction programming language

Key strengths:

- software infrastructure
- resource-constrained applications”

C++ 11

The Future is here

-- <http://www.infoq.com/presentations/Cplusplus-11-Bjarne-Stroustrup>

Linguagem + Biblioteca

C++ 14: <https://github.com/cplusplus/draft>

Document Number: N3797
Date: 2013-10-13
Revises: N3691
Reply to: Stefanus Du Toit
cxxeditor@gmail.com

Working Draft, Standard for Programming
Language C++

acesso em: 25 de Maio de 2014

C++ 11: ISO/IEC 14882:2011

Abstract

ISO/IEC 14882:2011 specifies requirements for implementations of the C++ programming language. The first such requirement is that they implement the language, and so ISO/IEC 14882:2011 also defines C++. Other requirements and relaxations of the first requirement appear at various places within ISO/IEC 14882:2011.

C++ is a general purpose programming language based on the C programming language as specified in ISO/IEC 9899:1999. In addition to the facilities provided by C, C++ provides additional data types, classes, templates, exceptions, namespaces, operator overloading, function name overloading, references, free store management operators, and additional library facilities.

Table 4 — Keywords

alignas	continue	friend	register	true
alignof	decltype	goto	reinterpret_cast	try
asm	default	if	return	typedef
auto	delete	inline	short	typeid
bool	do	int	signed	typename
break	double	long	sizeof	union
case	dynamic_cast	mutable	static	unsigned
catch	else	namespace	static_assert	using
char	enum	new	static_cast	virtual
char16_t	explicit	noexcept	struct	void
char32_t	export	nullptr	switch	volatile
class	extern	operator	template	wchar_t
const	false	private	this	while
constexpr	float	protected	thread_local	
const_cast	for	public	throw	

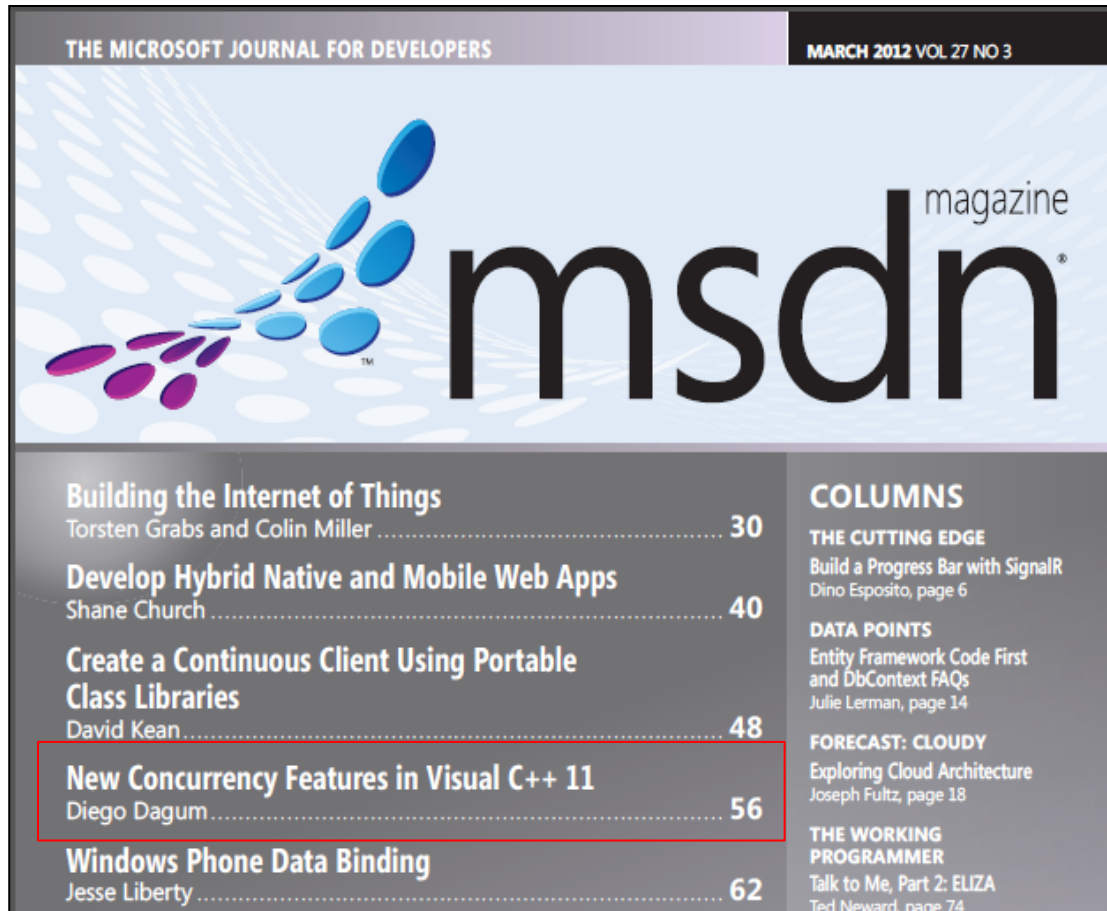
2.12 Keywords [lex.key] [2. Lexical conventions]

Compiladores



Summary of C++11 Feature Availability: <http://www.aristeia.com/C++11/C++11FeatureAvailability.htm>

Programação Concorrente com C++ 11



artigo online: <http://msdn.microsoft.com/en-us/magazine/hh852594.aspx>
código fonte: <http://archive.msdn.microsoft.com/mag201203CPP>

This article discusses Visual C++ 11, a prerelease technology. All related information is subject to change.

This article discusses:

- Parallel execution
- Asynchronous tasks
- Threads
- Variables and exceptions
- Synchronization
- Atomic types
- Mutexes and locks
- Condition variables

Technologies discussed:

Visual C++ 11

Code download available at:

code.msdn.microsoft.com/mag201203CPP

DIEGO DAGUM is a software developer with more than 20 years of experience. He's currently a Visual C++ community program manager with Microsoft.

THANKS to the following technical experts for reviewing this article:
David Cravey, Alon Flies, **Fabio Galuppo** and Marc Gregoire

Multi-threading with `std::`

library

Multi-threading

Atomic and thread support
Support for atomics and threads:


● **Headers**

<code><atomic></code>	Atomic (header)
<code><thread></code>	Thread (header)
<code><mutex></code>	Mutex (header)
<code><condition_variable></code>	Condition variable (header)
<code><future></code>	Future (header)

<http://www.cplusplus.com/reference/multithreading/>


#include <thread>

header


<thread> 

Thread

Header that declares the `thread` class and the `this_thread` namespace:

 Classes

<code>thread</code>	Thread (class)
---------------------	-----------------

 Classes

<code>this_thread</code>	This thread (namespace)
--------------------------	--------------------------

<http://www.cplusplus.com/reference/thread/>


```

void thread_func()
{
    std::stringstream ss;
    ss << "inside thread with id "<< std::this_thread::get_id() << "..." << std::endl;
    std::cout << ss.str();
}

std::thread t1(thread_func);
std::thread t2{ thread_func() };

t1.join();

t2.join();
//if (t2.joinable()) t2.detach();
try
{
    t2.join();
}
catch (std::system_error& e)
{
    auto err = e.code();
    if (err.value() == (int)std::errc::invalid_argument)
        std::cout << "The thread object is not joinable" << std::endl;
    else
        std::cout << e.what() << std::endl;
}

```



```

template <class RandomAccessIterator>
void merge_sort(RandomAccessIterator first, RandomAccessIterator end)
{
    if (end - first <= 1) return;

    auto mid = std::distance(first, end) / 2;
    merge_sort(first, first + mid);
    merge_sort(first + mid, end);
    merge(first, end, mid);
}

#include <sstream>

template <class RandomAccessIterator>
void parallel_merge_sort(RandomAccessIterator first, RandomAccessIterator end)
{
    if (end - first <= 1) return;

    auto size = std::distance(first, end);
    auto nothreads = std::thread::hardware_concurrency();

    std::vector<std::thread> threads;
    auto it = first;

    auto f = [](RandomAccessIterator f, RandomAccessIterator e)
    {
        std::stringstream ss;
        ss << std::distance(f, e) << " " << &f << " " << &e << std::endl;
        std::cout << ss.str();
        merge_sort(f, e);
    };

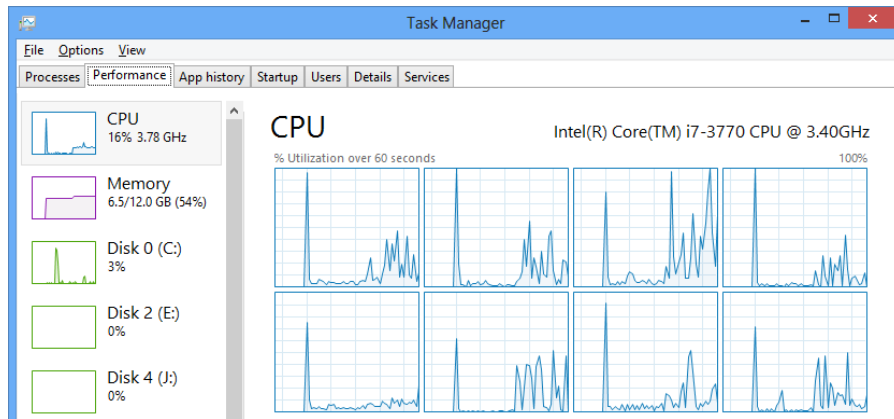
    auto stride = static_cast<unsigned>(size * (1.0 / nothreads));
    for (unsigned i = 1; i < nothreads; ++i, it += stride)
    {
        std::thread t(f, it, it + stride);
        threads.push_back(std::move(t));
    }
    std::thread t(f, it, end);
    threads.push_back(std::move(t));
    for (auto& th : threads)
        th.join();

    merge(first, end, std::distance(first, end) / 2);
}

```

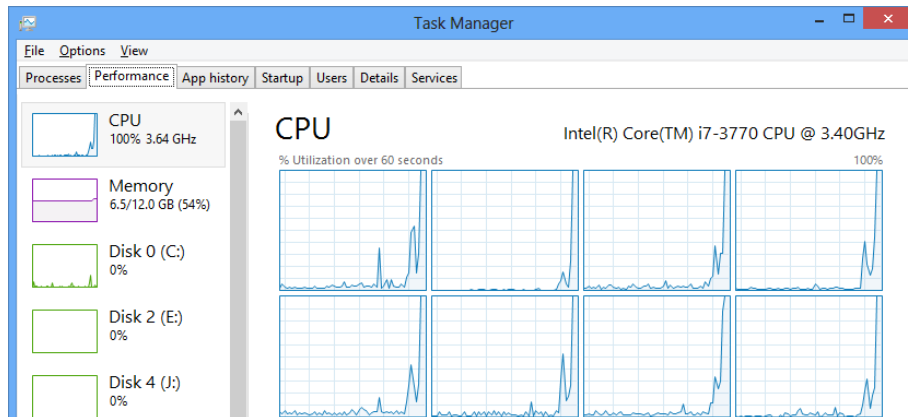
Embarrassingly Parallel

Serial



```
C:\WINDOWS\system32\cmd.exe
container size = 100000000
737043 589503 836958 958228 652229 238270 968502 971339 228045 293907
589143 819903 809199 447004 960755 856415 47386 643549 991738 912776
1 1 1 1 1 1 1 1 1 1
1000000 1000000 1000000 1000000 1000000 1000000 1000000 1000000 1000000 1000000
22215 ms
Press any key to continue . . .
```

Paralelo



```
C:\WINDOWS\system32\cmd.exe
container size = 100000000
223760 770230 401302 323600 405634 618706 638433 760310 670150 309673
79729 179867 303650 858692 520350 133303 904469 269152 200469 239378
12500000 00E1F9CC 00E1F9D0
12500000 00F1F7A4 00F1F7A8
12500000 0101F7AC 0101F7B0
12500000 18F4F9D0 18F4F9D4
12500000 1904FC20 1904FC24
12500000 1918F820 1918F824
12500000 192CFC88 192CFC8C
12500000 1940FCCC 1940FCD0
1 1 1 1 1 1 1 1 1 1
1000000 1000000 1000000 1000000 1000000 1000000 1000000 1000000 1000000 1000000
4595 ms
Press any key to continue . . .
```

#include <mutex>

header

<mutex>

Mutex

Header with facilities that allow *mutual exclusion* (mutex) of concurrent execution of critical sections of code, allowing to explicitly avoid data races.

It contains *mutex types*, *lock types* and specific *functions*:

- **Mutex types** are *lockable types* used to protect access to a *critical section* of code: *locking* a mutex prevents other threads from locking it (exclusive access) until it is *unlocked*: `mutex`, `recursive_mutex`, `timed_mutex`, `recursive_timed_mutex`.
- **Locks** are objects that manage a mutex by associating its access to their own lifetime: `lock_guard`, `unique_lock`.
- **Functions** to lock multiple mutexes simultaneously (`try_lock`, `lock`) and to directly prevent concurrent execution of a specific function (`call_once`).

Classes

Mutexes

<code>mutex</code>	Mutex class (class)
<code>recursive_mutex</code>	Recursive mutex class (class)
<code>timed_mutex</code>	Timed mutex class (class)
<code>recursive_timed_mutex</code>	Recursive timed mutex (class)

Locks

<code>lock_guard</code>	Lock guard (class template)
<code>unique_lock</code>	Unique lock (class template)

Other types





<code>once_flag</code>	Flag argument type for <code>call_once</code> (class)
<code>adopt_lock_t</code>	Type of <code>adopt_lock</code> (class)
<code>defer_lock_t</code>	Type of <code>defer_lock</code> (class)
<code>try_to_lock_t</code>	Type of <code>try_to_lock</code> (class)

Functions

<code>try_lock</code>	Try to lock multiple mutexes (function template)
<code>lock</code>	Lock multiple mutexes (function template)
<code>call_once</code>	Call function once (public member function)

<http://www.cplusplus.com/reference/mutex/>

#include <condition_variable>

header	
<condition_variable> 	
Condition variable	
Header that declares <i>condition variable</i> types:	
 Classes	
condition_variable	Condition variable (class)
condition_variable_any	Condition variable (any lock) (class)
 Enum classes	
cv_status	Condition variable status (enum class)
 Functions	
notify_all_at_thread_exit	Notify all at thread exit (function)

http://www.cplusplus.com/reference/condition_variable/

```

#include <mutex>
#include <condition_variable>
#include <queue>

template<typename T>
struct blocking_queue final
{
    void enqueue(const T& item)
    {
        std::lock_guard<std::mutex> lock(CR_);
        Q_.push(item);
        NonEmpty_.notify_one();
    }

    T dequeue()
    {
        std::unique_lock<std::mutex> lock(CR_);
        NonEmpty_.wait(lock, [this]() { return !Q_.empty(); });
        //or
        //while (Q_.empty()) NonEmpty_.wait(lock);
        T temp = Q_.front();
        Q_.pop();
        return temp;
    }

private:
    std::condition_variable NonEmpty_;
    std::mutex CR_;
    std::queue<T> Q_;
};

blocking_queue<int> q;

//consumer
auto deq_f = std::async([&q]() {
    int value;
    while ((value = q.dequeue()) != 0)
        std::cout << value << " ";
});

//producer
std::default_random_engine engine(static_cast<unsigned>(std::time(nullptr)));
std::uniform_int_distribution<int> rnd(1, 100000);


for (int i = 0; i < 10; ++i)
{
    for (int j = 0; j < 10; ++j) q.enqueue(rnd(engine));
    std::this_thread::sleep_for(std::chrono::seconds(1));
}

std::this_thread::sleep_for(std::chrono::seconds(3));
q.enqueue(0);

```

#include <future>


header

<future> 

Future

Header with facilities that allow asynchronous access to values set by specific providers, possibly in a different thread.

Each of these providers (which are either `promise` or `packaged_task` objects, or calls to `async`) share access to a *shared state* with a `future` object: the point where the *provider* makes the *shared state* ready is synchronized with the point the `future` object accesses the *shared state*.

 **Classes**

Providers


<code>promise</code>	Promise (class template)
<code>packaged_task</code>	Packaged task (class template)

Futures

<code>future</code>	Future (class template)
<code>shared_future</code>	Shared future (class template)

Other types

<code>future_error</code>	Future error exception (class)
<code>future_errc</code>	Error conditions for future objects (enum class)
<code>future_status</code>	Return value for timed future operations (enum class)
<code>launch</code>	Launching policy for <code>async</code> (enum class)

 **Functions**

Providers

<code>async</code>	Call function asynchronously (function template)
--------------------	---

Other functions

<code>future_category</code>	Return future category (function)
------------------------------	------------------------------------

<http://www.cplusplus.com/reference/future/>

```
//future and async
std::future<int> f =
    std::async([]() -> int {
        std::this_thread::sleep_for(std::chrono::seconds(2));
        return 100;
    });

std::cout << f.get() << std::endl;
```

```
auto g = [](std::promise<int>& p, int value) {
    if (value <= 0)
    {
        auto e = std::make_exception_ptr(std::runtime_error("value less
        p.set_exception(e);
        return;
    }

    std::this_thread::sleep_for(std::chrono::seconds(2));
    p.set_value(value * 100);
};

//promise
std::promise<int> p1, p2;
std::thread t1(g, std::ref(p1), 100);
std::thread t2(g, std::ref(p2), -1);

auto f1 = p1.get_future();
auto f2 = p2.get_future();

f1.wait();
f2.wait();

std::cout << f1.get() << std::endl;
try
{
    std::cout << f2.get() << std::endl;
}
catch (const std::exception& e)
{
    std::cout << "exception caught: " << e.what() << std::endl;
}

t1.join();
t2.join();
```

```
//packaged_task
auto h = [](int a, int b)
{
    std::this_thread::sleep_for(std::chrono::seconds(2));
    int result = a * b;
    if (result <= 0)
        throw std::runtime_error("result less than or equal to 0");
    return result;
};

std::packaged_task<int(int, int)> pt1(h);
std::future<int> f3 = pt1.get_future();
std::thread t3(std::move(pt1), 10, 90);

std::packaged_task<int(int, int)> pt2(h);
std::future<int> f4 = pt2.get_future();
std::thread t4(std::move(pt2), 10, -90);
t3.detach();
t4.detach();

std::cout << f3.get() << std::endl;
try
{
    std::cout << f4.get() << std::endl;
}
catch (const std::exception& e)
{
    std::cout << "exception caught: " << e.what() << std::endl;
}
```


Wrapping up

```
//TSP -> NN -> Generations( g, ForkJoin ( n, SA -> 2-OPT ) ) -> TSP'
void Pipeline1(tsp_class& tsp_instance, unsigned int number_of_tasks, unsigned int number_of_generations)
{
    #pragma region "PipelineConfiguration"
    auto a = Args<General_args_type>(make_General_args(number_of_generations, number_of_tasks));
    auto sa = Args<SA_args_type>(make_SA_args(1000.0, 0.00001, 0.999, 400));
    auto aco = Args<ACO_args_type>();
    auto ga = Args<GA_args_type>();



    const char* pipeline_description = "TSP -> NN -> Generations( g, ForkJoin ( n, SA -> 2-OPT ) ) -> TSP'";
    display_args(pipeline_description, a, sa, aco, ga);

    auto g = a[0].number_of_iterations_or_generations;
    auto n = a[0].number_of_tasks_in_parallel;
    auto _TSP = TSP(just(tsp_instance));
    auto _DisplayInput = Display("TSP INPUT", DisplayFlags::All);
    auto _NN = Measure(NN(), Display("NEAREST NEIGHBOUR", DisplayFlags::EmitMathematicaGraphPlot));
    auto _SA_2OPT = Chain(SA(sa[0].initial_temperature, sa[0].stopping_criteria_temperature,
        sa[0].decreasing_factor, sa[0].monte_carlo_steps), _2OPT());
    auto _ForkJoin = [](unsigned int n, TSP::transformer_type map_fun){ return Measure(ForkJoin(n, map_fun)); };
    auto _DisplayOutput = Display("TSP OUTPUT", DisplayFlags::EmitMathematicaGraphPlot);
    #pragma endregion

    //TSP -> NN -> Generations( g, ForkJoin ( n, SA -> 2-OPT ) ) -> TSP'
    auto result = _TSP
        .map(_DisplayInput)
        .map(_NN)
        .map(Generations(g, _ForkJoin(n, _SA_2OPT)))
        .map(_DisplayOutput);
}
```

<http://bit.ly/1hoODIh>

Task Manager						
File	Options	View				
Processes	Performance	App history	Startup	Users	Details	Services
Name	PID	Status	User name	CPU	Memory	Private mem...
TCSP.exe	1920	Running	Fabio Galuppo	96	2,382 K	43



<http://bit.ly/1ppSmKN>

```
struct ForkJoin
{
    typedef TSP::transformer_type F;

    ForkJoin(unsigned int number_of_tasks_in_parallel, F f, bool propagate_inferior_results = false) :
        NumberOfTasksInParallel_(number_of_tasks_in_parallel),
        Func_(f),
        PropagateInferiorResults_(propagate_inferior_results)
    {
    }

    TSP operator()(TSP::T t)
    {
        if (0 == NumberOfTasksInParallel_)
            return TSP(t);

        std::vector<std::future<tsp_async_state>> tsp_solutions;


        std::cout << "START SOLUTION:" << std::endl;

        auto tsp_instance = ref(t);
        display_solution(tsp_instance);

        for (unsigned int i = 0; i < NumberOfTasksInParallel_; ++i)
        {
            //async
            tsp_solutions.push_back(std::async(std::launch::async, [&, i, tsp_instance]
            {
                Maybe a = just(tsp_instance);
                std::unique_ptr<ForkJoin_args> args(new ForkJoin_args(i, NumberOfTasksInParallel_));
                a.second = args.get();
                return tsp_async_state(Func_(a));
            }));
        }

        std::cout << "CANDIDATE SOLUTIONS:" << std::endl;
        std::vector<tsp_class> candidate_solutions;
        int i = 0;
        for (auto& tsp_solution : tsp_solutions)
        {
            //async result
            auto candidate = tsp_solution.get();
            if (has(candidate.state))
            {
                candidate_solutions.push_back(ref(candidate.state));
                display(ref(candidate.state), false);
            }
        }
    }
}
```

Intel Threading Building Blocks (TBB)



Intel® Threading Building Blocks
C and C++ template library for creating high performance, scalable parallel applications

Included in Intel®
Parallel Studio XE &
Intel® Cluster Studio

Generic Parallel Algorithms	Concurrent Containers	Task Scheduler	Synchronization Primitives	Memory Allocation	Miscellaneous
<code>parallel_for(range)</code> <code>parallel_reduce</code> <code>parallel_for_each(begin, end)</code> <code>parallel_do</code> <code>parallel_invoke</code> <code>pipeline</code> <code>parallel_pipeline</code> <code>parallel_sort</code> <code>parallel_scan</code> <code>flow::graph</code> <code>parallel_deterministic_reduce</code>	<code>concurrent_hash_map</code> <code>concurrent_queue</code> <code>concurrent_bounded_queue</code> <code>concurrent_vector</code> <code>concurrent_unordered_map</code> <code>concurrent_priority_queue</code> <code>concurrent_unordered_set</code>	<code>task</code> <code>task_group</code> <code>structured_task_group</code> <code>task_group_context</code> <code>task_scheduler_init</code> <code>task_scheduler_observer</code>	<code>atomic</code> <code>mutex</code> <code>recursive_mutex</code> <code>spin_mutex</code> <code>spin_rw_mutex</code> <code>queuing_mutex</code> <code>queuing_rw_mutex</code> <code>reader_writer_lock</code> <code>critical_section</code> <code>condition_variable</code> <code>null_mutex</code> <code>null_rw_mutex</code>	<code>tbb_allocator</code> <code>cache_aligned_allocator</code> <code>scalable_allocator</code> <code>zero_allocator</code> <code>memory_pool</code>	<code>thread</code> <code>tick_count</code> <code>captured_exception</code> <code>moveable_exception</code> <code>enumerable_thread_specific</code> <code>combinable</code>



Optimized Threading Functions
running on Windows*, Linux*, OS X* & more



<https://software.intel.com/en-us/intel-tbb>

<https://www.threadingbuildingblocks.org/>

Visual C++ ConcRT e Parallel Patterns Library (PPL)

PPL Compatibility

Intel® Threading Building Blocks (Intel® TBB) 2.2 introduces features based on joint discussions between the Microsoft Corporation and Intel Corporation. The features establish some degree of compatibility between Intel® TBB and Microsoft Parallel Patterns Library (PPL) development software.

Each feature appears in namespace `tbb`. Each feature can be injected into namespace `Concurrency` by including the file `"tbb/compat/pp1.h"`. Following is the list of features:

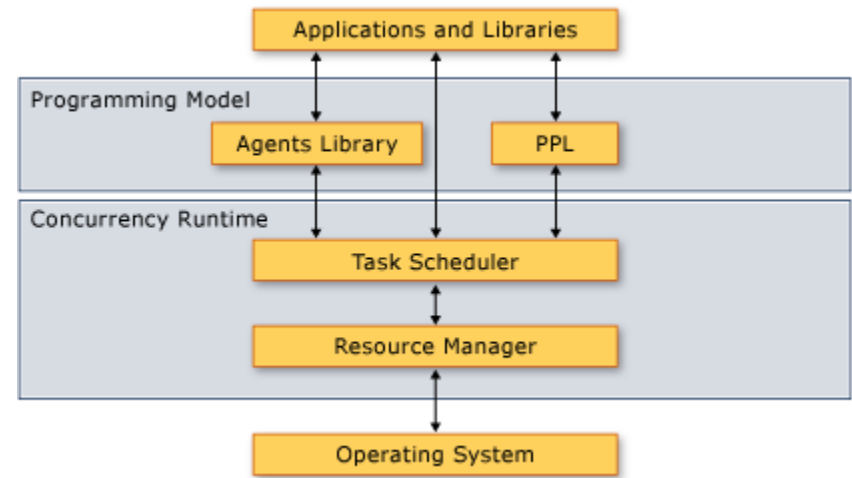
The following table lists the features and provides links to additional information.

Feature	Link
<code>parallel_for(first,last,f)</code>	parallel_for Template Function
<code>parallel_for(first,last,step,f)</code>	parallel_for Template Function
<code>parallel_for_each</code>	parallel_for_each Template Function
<code>parallel_invoke</code>	parallel_invoke Template Function
<code>critical_section</code>	critical_section
<code>reader_writer_lock</code>	reader_writer_lock Class
<code>task_handle</code>	task_handle Template Class
<code>task_group_status</code>	task_group_status Enum
<code>task_group</code>	task_group Class
<code>make_task</code>	make_task Template Function
<code>structured_task_group</code>	structured_task_group Class
<code>is_current_task_group_cancelling</code>	is_current_task_group_canceling Function
<code>improper_lock</code>	Specific Exceptions
<code>invalid_multiple_scheduling</code>	Specific Exceptions
<code>missing_wait</code>	Specific Exceptions

For `parallel_for`, only the variants listed in the table are injected into namespace `Concurrency`.

CAUTION

Because of different environments and evolving specifications, the behavior of the features can differ between the Intel® TBB and PPL implementations.



<http://msdn.microsoft.com/en-us/library/ee207192.aspx>

The PPL provides the following features:

- *Task Parallelism*: a mechanism to execute several work items (tasks) in parallel
- *Parallel algorithms*: generic algorithms that act on collections of data in parallel
- *Parallel containers and objects*: generic container types that provide safe concurrent access to their elements

<http://msdn.microsoft.com/en-us/library/dd492418.aspx>


Introdução a *Multi-threading* com C++ 11

(C++ 11 == ISO/IEC 14882:2011)

Fabio Galuppo, M.Sc.

<http://fabiogaluppo.com>

fabiogaluppo@acm.org



First year awarded:
2002

Number of MVP Awards:
11

Technical Expertise:
Visual C++

Technical Interests:
Visual C#, Visual F#